Graph-Based Polypharmacy Risk Networks for Adverse Event Prediction in ICU Patients: A MIMIC-IV Analysis

Author Name
Affiliation
Email

Abstract—**Polypharmacy – the concurrent use of multiple medications – is common in intensive care and is a recognized driver of adverse events, yet most machine-learning studies on MIMIC-IV target mortality prediction rather than modeling drug-interaction risk directly. We address this gap by constructing a population-level drug co-administration graph from 9,974 adult ICU stays and comparing three patient-level feature representations – node2vec graph embeddings, hand-crafted graph-theoretic statistics, and a one-hot tabular baseline – for predicting three polypharmacy-relevant outcomes: acute kidney injury (AKI, KDIGO-based), a delirium proxy, and a bleeding proxy. Across four classifiers and stratified 5-fold cross-validation, we find a mixed but interpretable pattern: the tabular baseline wins when given unrestricted access to drug identity (AKI AUROC 0.763, bleeding AUROC 0.864), while graph embeddings outperform a leakage-corrected tabular baseline for delirium (AUROC 0.876 vs. 0.858, paired bootstrap Δ=+0.018, 95% CI [0.011, 0.025], $p < 0.0001$), a result that replicates on an independent, zero-overlap second cohort subsample (Δ=+0.015, 95% CI [0.008, 0.022]). We further report a label-leakage correction identified during model development (delirium label circularity with tabular drug features, corrected AUROC $0.944 \rightarrow 0.858$) and perturbation-based, reliability-filtered drug-risk rankings per outcome. We present this comparison as evidence for when graph-based representations add value over tabular baselines in EHR-based polypharmacy modeling.**

Index Terms—**polypharmacy, graph embeddings, node2vec, MIMIC-IV, electronic health records, acute kidney injury, delirium, adverse drug events**

I. Introduction

Machine learning research on the MIMIC-IV critical care database has concentrated heavily on in-hospital mortality prediction, leaving polypharmacy-specific adverse event modeling comparatively underexplored. Polypharmacy – typically defined as the concurrent administration of multiple medications during a single stay – is near-universal in the ICU and is mechanistically linked to adverse drug events including nephrotoxicity, delirium, and bleeding complications, yet the *combinatorial* structure of drug co-administration is rarely modeled explicitly; most prior EHR risk models treat medication exposure as an unordered feature bag rather than as a network of interacting entities.

This paper asks a narrower, more falsifiable question than "can graphs model polypharmacy risk": under what conditions does a graph-based representation of drug co-administration outperform a conventional tabular representation with the same underlying information? We build a population-level co-administration graph over generic drug names, learn node2vec [1] embeddings for each drug, and compare three patient-level feature representations head-to-head, on identical cohorts, labels, classifiers, and cross-validation protocol, for three outcomes: acute kidney injury (AKI), a delirium proxy, and a bleeding proxy.

Our contribution is threefold: (1) a three-way feature-representation comparison (graph embedding vs. graph-theoretic vs. tabular) evaluated under matched conditions across three clinically distinct outcomes; (2) a documented and corrected label-leakage circularity between the delirium outcome definition and tabular drug features, together with a paired significance test validating the corrected result; and (3) a confirmatory replication of the headline result on an independent cohort subsample, directly addressing the generalizability concern that a single-subsample result invites.

II. Related Work

MIMIC-based clinical prediction. MIMIC-IV [2] has enabled a large body of work on ICU outcome prediction, predominantly targeting mortality and length-of-stay as the outcome of interest; this concentration is the specific gap this paper targets by modeling polypharmacy-relevant adverse events instead.

Graph-based EHR modeling. Representing patients or clinical entities as graphs – and learning representations over them via methods such as node2vec [1] – has been explored for diagnosis and comorbidity modeling; we apply this framework specifically to drug co-administration structure, which has received comparatively less direct empirical comparison against tabular baselines under matched evaluation protocol.

Interpretability. We use SHAP [3] for the tabular baseline's tree-based models, and a perturbation (leave-one-drug-out) analysis for the graph-embedding models, since raw embedding dimensions are not individually interpretable.

III. Methods

A. Cohort construction

We extracted adult patients (age ≥18) with a first ICU stay of length ≥24 hours from MIMIC-IV via Google BigQuery, yielding 68,013 eligible admissions. For BigQuery cost and local compute tractability, we drew a fixed-seed random

TABLE I
COHORT ATTRITION.

Filtering step	n (admissions)
Adult patients, first ICU stay per admission	85,242
ICU LOS ¿= 24h	68,013
Random subsample (seed=42, compute/cost constraint)	10,000
¿= 3 distinct generic drugs administered	9,974
Final cohort (all labels attached)	9,974

subsample of 10,000 admissions (seed=42, ≈14.7% of the eligible cohort), then excluded admissions with fewer than 3 distinct administered drugs (no meaningful polypharmacy signal), yielding a final analytic cohort of $n = 9{,}974$. Table I reports the full filtering flow.

This subsampling is a stated scope limitation, not a hidden one: to test whether headline findings generalize beyond the specific subsample drawn, we separately replicate our primary result on an independent, zero-overlap second 10,000-admission subsample (Section IV-C).

B. Outcome labels

We define three binary outcomes, each an explicit, documented proxy rather than an adjudicated clinical diagnosis:

- **AKI**: KDIGO stage ≥ 1 [4], computed via MIMIC-IV's derived `kdigo_stages` table. Positive in 7,312/9,974 (73.3%) of the cohort.
- **Delirium (proxy)**: administration of an antipsychotic or benzodiazepine commonly used for agitation (haloperidol, quetiapine, olanzapine, risperidone, ziprasidone, lorazepam, midazolam, diazepam) **AND** a positive CAM-ICU [5] chartevent (2 of the 4 CAM-ICU criteria are directly captured in MIMIC-IV chartevents: altered mental status and inattention; the full 4-criterion clinical algorithm is not reconstructed). Positive in 2,679/9,974 (26.9%).
- **Bleeding (proxy)**: a new packed red blood cell (PRBC) transfusion order after ICU day 2. Positive in 876/9,974 (8.8%).

All three are stated limitations in Section V-A: they are defensible, literature-grounded proxies, not gold-standard adjudicated outcomes.

C. Drug co-administration graph

Prescription records were normalized to approximate generic names via a documented, deterministic cleaning pipeline (dose/form-suffix stripping plus a small brand-to-generic synonym dictionary). For each admission, two drugs are considered co-administered if given on an overlapping calendar day. The *population-level* co-administration graph aggregates this across the full cohort: nodes are distinct generic drug names, and edge weight is the number of distinct admissions in which the pair was co-administered. Table II reports summary statistics; Figure 1 visualizes the top-60 highest-degree drugs, and Figure 2 shows the full degree distribution.

TABLE II
POPULATION CO-ADMINISTRATION GRAPH SUMMARY STATISTICS.

Statistic	Value
Nodes (distinct drugs)	2,091
Edges (co-administered pairs)	160,276
Density	0.0733
Avg. clustering coefficient	0.0011
Connected components	5
Mean degree	153.3
Median degree	49.0
Max degree	1573

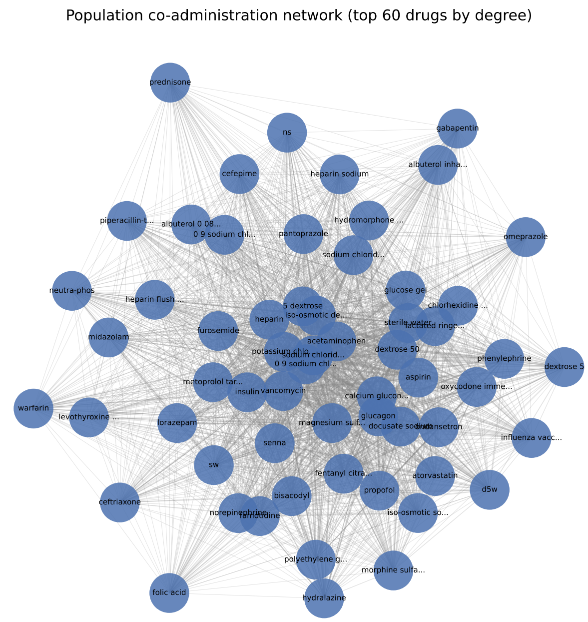

Fig. 1. Population co-administration network, top 60 drugs by degree. Node size scales with degree; edge width scales with co-administration weight.

D. Feature representations

We compare three patient-level feature representations, computed for every admission:

1) **Graph embedding**: node2vec [1] is trained on the weighted population graph (walk length 40, 20 walks per node, embedding dimension 64, window 10, default $p{=}1, q{=}1$ – not hyperparameter-tuned, a deliberate choice for a fair, fixed-hyperparameter comparison across all three feature sets, discussed further in Section V). Each patient's feature vector is the mean embedding of their administered drugs.
2) **Graph-theoretic**: sum and max pairwise edge weight among a patient's drugs, and the density of their drug-induced subgraph.
3) **Tabular baseline**: one-hot indicators for every administered drug, plus age, sex, and a comorbidity count (distinct ICD diagnosis codes).

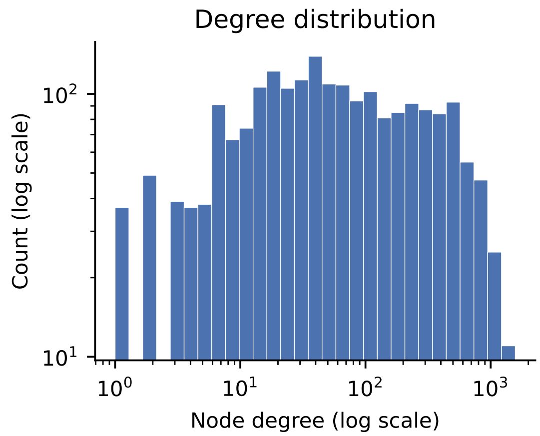

Fig. 2. Degree distribution of the population co-administration graph (log-log).

E. Modeling and evaluation

For every outcome × feature-set combination, we train four classifiers with fixed (untuned) hyperparameters – logistic regression, random forest (300 trees), XGBoost [6] (300 trees, depth 4), and a 2-layer MLP (64, 32 units) – under stratified 5-fold cross-validation. We report AUROC, AUPRC, F1, and Brier score, with AUROC 95% confidence intervals from 1,000 bootstrap resamples of the pooled out-of-fold predictions.

F. Label-leakage correction

The delirium label is partly *defined* by administration of specific drugs (Section III-B); the tabular baseline's one-hot features include columns for those same drugs, so a tabular model can trivially reconstruct much of the label via this circularity rather than genuine predictive signal. We identified this after observing an implausibly high tabular AUROC (0.944, XGBoost) for delirium. We excluded the delirium-defining drug columns from the tabular feature set *for the delirium outcome only* (AKI and bleeding labels are not drug-defined and are unaffected) and refit all four tabular models, yielding a corrected AUROC of 0.858 (XGBoost).

IV. Results

A. Three-way feature comparison

Table III reports the full outcome × feature-set × model grid; Figure 3 summarizes the best-model AUROC per feature set and outcome. The comparison is mixed rather than a uniform win for any single representation: the tabular baseline achieves the highest AUROC for AKI (0.763, XGBoost) and bleeding (0.864, XGBoost), while the graph embedding achieves the highest AUROC for delirium (0.876, XGBoost) – exceeding even the leakage-corrected tabular result (0.858).

We caution that AUPRC for AKI (base rate 73.3%) should be interpreted relative to that high base rate rather than in isolation; AUROC is the more informative metric for this outcome, and we lead with it throughout.

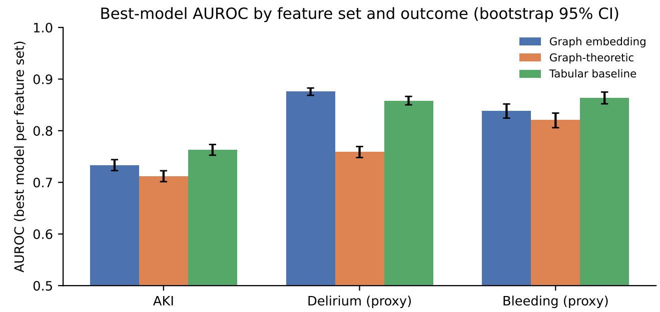

Fig. 3. Best-model AUROC by feature set and outcome, with bootstrap 95% CI error bars.

B. Statistical significance of the delirium result

Because the embedding-vs-tabular delirium comparison is our headline claim, we tested it with a paired bootstrap on *identical* cross-validation fold assignments for both feature sets (rather than comparing two independently-computed confidence intervals), so the significance estimate reflects a true paired comparison. Table IV reports the result: the embedding's AUROC advantage over the leak-free tabular baseline is statistically significant (Δ=+0.0178, 95% CI [0.0112, 0.0249], $p < 0.0001$).

C. Confirmatory replication

To address whether this result is specific to the particular 10,000-admission subsample drawn, we repeated the entire pipeline – cohort draw, graph construction, node2vec training, and the paired significance test – on a second, independent subsample with *zero admission overlap* with the original sample (seed=123, n = 9,976 after filtering). The effect direction and approximate magnitude replicated (Table V): Δ=+0.0152, 95% CI [0.0081, 0.0218], overlapping the original sample's confidence interval.

D. Interpretability

For each outcome's best embedding model, we identify risk-associated drugs via perturbation analysis: for high-risk patients (top decile of predicted probability), we remove each drug from their exposure set, re-score, and average the resulting prediction delta across patients. To avoid reporting single-patient noise as signal, we exclude any drug supported by fewer than 10 high-risk patients before ranking (Figure 4). These rankings have *not* been cross-referenced against a clinical drug-drug-interaction database and should not be interpreted as validated interactions without that manual check.

V. Discussion

The three-way comparison does not support a blanket claim that graph embeddings outperform tabular baselines for polypharmacy risk modeling. Instead, it supports a more specific and, we think, more useful claim: a tabular representation with unrestricted access to raw drug identity has a strict informational advantage and wins when that access is available

TABLE III
AUROC (95% bootstrap CI), AUPRC, F1, and Brier score for every outcome × feature-set × model combination. delirium_label / tabular rows are the leakage-corrected refit (see Section III).

Outcome	Feature set	Model	AUROC	95% CI	AUPRC	F1	Brier
AKI	Graph embedding	random forest	0.733	[0.723, 0.744]	0.880	0.851	0.170
		xgboost	0.726	[0.716, 0.738]	0.878	0.845	0.172
		logistic regression	0.720	[0.709, 0.732]	0.869	0.846	0.172
		mlp	0.631	[0.618, 0.643]	0.806	0.778	0.296
AKI	Graph-theoretic	mlp	0.712	[0.701, 0.722]	0.866	0.842	0.175
		logistic regression	0.709	[0.698, 0.720]	0.863	0.843	0.176
		xgboost	0.706	[0.695, 0.716]	0.864	0.837	0.177
		random forest	0.662	[0.650, 0.673]	0.841	0.817	0.193
AKI	Tabular baseline	xgboost	0.763	[0.753, 0.773]	0.896	0.851	0.162
		random forest	0.759	[0.748, 0.770]	0.890	0.853	0.164
		logistic regression	0.686	[0.674, 0.697]	0.848	0.808	0.210
		mlp	0.659	[0.647, 0.670]	0.841	0.779	0.294
Delirium (proxy)	Graph embedding	xgboost	0.876	[0.869, 0.883]	0.735	0.625	0.121
		random forest	0.849	[0.841, 0.858]	0.710	0.569	0.133
		mlp	0.849	[0.842, 0.857]	0.664	0.621	0.182
		logistic regression	0.849	[0.841, 0.856]	0.659	0.587	0.135
Delirium (proxy)	Graph-theoretic	mlp	0.759	[0.748, 0.769]	0.561	0.433	0.161
		xgboost	0.755	[0.744, 0.765]	0.551	0.419	0.162
		logistic regression	0.738	[0.727, 0.749]	0.521	0.349	0.167
		random forest	0.718	[0.707, 0.729]	0.505	0.438	0.176
Delirium (proxy)	Tabular baseline	xgboost	0.858	[0.850, 0.866]	0.704	0.606	0.128
		random forest	0.848	[0.839, 0.856]	0.693	0.540	0.135
		logistic regression	0.784	[0.773, 0.795]	0.557	0.565	0.173
		mlp	0.772	[0.762, 0.782]	0.566	0.545	0.208
Bleeding (proxy)	Graph embedding	xgboost	0.838	[0.824, 0.852]	0.384	0.237	0.065
		logistic regression	0.810	[0.795, 0.825]	0.317	0.154	0.069
		random forest	0.787	[0.770, 0.803]	0.337	0.143	0.069
		mlp	0.783	[0.768, 0.798]	0.260	0.312	0.101
Bleeding (proxy)	Graph-theoretic	mlp	0.821	[0.806, 0.834]	0.320	0.109	0.069
		logistic regression	0.817	[0.802, 0.831]	0.318	0.153	0.070
		xgboost	0.807	[0.792, 0.821]	0.298	0.138	0.070
		random forest	0.761	[0.743, 0.776]	0.250	0.195	0.077
Bleeding (proxy)	Tabular baseline	xgboost	0.864	[0.852, 0.875]	0.442	0.338	0.062
		random forest	0.858	[0.846, 0.870]	0.431	0.157	0.064
		mlp	0.764	[0.746, 0.779]	0.267	0.326	0.122
		logistic regression	0.722	[0.703, 0.742]	0.249	0.317	0.114

TABLE IV
Paired bootstrap significance test: node2vec embedding vs. leakage-corrected tabular baseline, delirium outcome, identical 5-fold CV splits, 2,000 resamples.

Quantity	Value
AUROC (embedding)	0.8760
AUROC (tabular, leak-free)	0.8582
Observed ΔAUROC	0.0178
95% CI on Δ	[0.0112, 0.0249]
p (approx., two-sided)	0.0000

(AKI, bleeding), but once that advantage is removed by a genuine label-leakage correction (delirium), the compressed 64-dimensional graph embedding recovers signal that the reduced tabular baseline does not – and this advantage is both statistically significant on identical folds and replicates on an independent cohort subsample.

We note that our node2vec hyperparameters (p=1, q=1, no walk-length or dimension search) were fixed, not tuned, specifically to keep the three-way comparison fair. This cuts

TABLE V
Confirmatory replication on an independent, zero-overlap second 10,000-admission subsample (graph and embeddings rebuilt from scratch).

Quantity	Original sample	Replication
ΔAUROC (emb. − tab.)	0.0178	0.0152
95% CI	[0.0112, 0.0249]	[0.0081, 0.0218]

in favor of the embedding result, not against it: an untuned graph embedding already outperforms a leakage-corrected tabular baseline, suggesting that hyperparameter search would plausibly widen rather than narrow this gap – a direction for future work rather than a caveat on the present one.

A. Limitations

(1) All three outcome labels are proxies, not adjudicated diagnoses – most notably the delirium proxy captures 2 of 4 CAM-ICU criteria. (2) The analytic cohort is a fixed-seed random subsample of the eligible MIMIC-IV population, not the full 68,013-admission eligible cohort; while our delirium

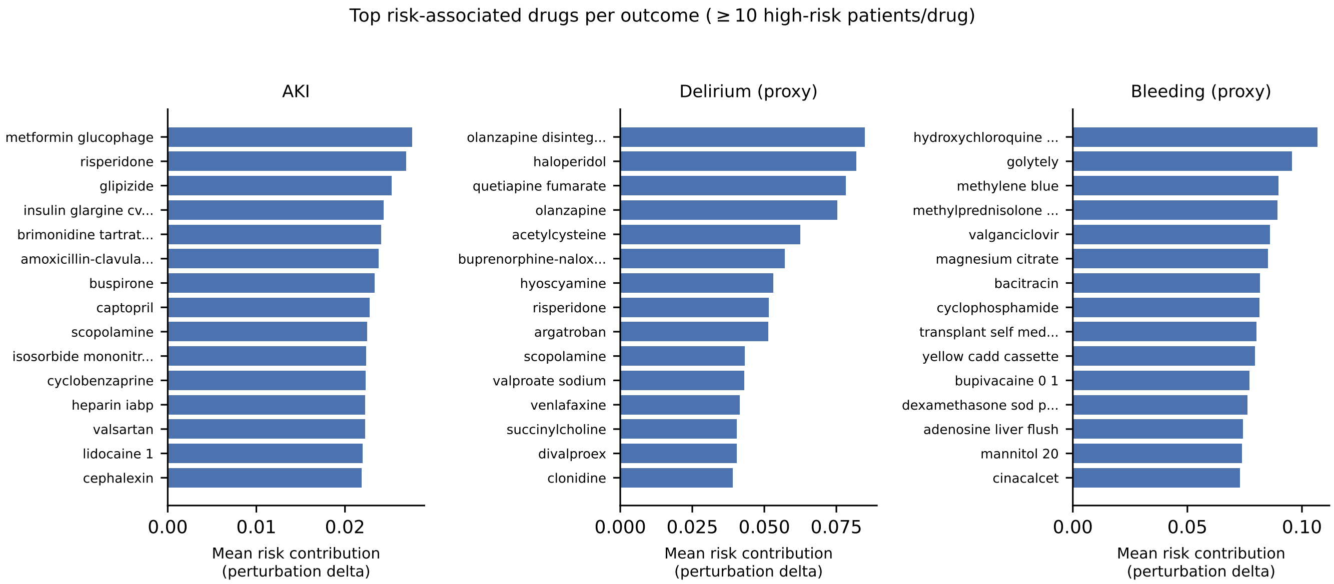

Fig. 4. Top perturbation-ranked risk-associated drugs per outcome (reliability-filtered, $\geq$10 high-risk patients per drug).

result replicates on an independent subsample, this was only tested for the headline claim, not for every outcome/feature-set combination. (3) Single-center, retrospective data (MIMIC-IV, derived from Beth Israel Deaconess Medical Center) with no external validation cohort. (4) No classifier hyperparameter tuning was performed, by design, for a fair fixed-hyperparameter comparison; absolute performance ceilings for any individual model are therefore likely understated. (5) Perturbation-based drug-risk rankings have not been manually cross-checked against a pharmacological drug-drug-interaction reference and should not be read as clinically validated without that step.

VI. Conclusion

We compared graph-embedding, graph-theoretic, and tabular feature representations for predicting three polypharmacy-relevant ICU adverse events on MIMIC-IV, under matched cohorts, labels, and evaluation protocol. The result is outcome-dependent: tabular features win when given unrestricted drug-identity access (AKI, bleeding), while graph embeddings outperform a leakage-corrected tabular baseline for delirium, a result that is statistically significant on paired folds and replicates on an independent cohort subsample. We further contribute a documented label-leakage correction and reliability-filtered, perturbation-based interpretability outputs. Future work includes node2vec hyperparameter search, external validation, and manual clinical cross-referencing of the identified risk-associated drug combinations.

References

[1] A. Grover and J. Leskovec, "node2vec: Scalable feature learning for networks," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 855–864.

[2] A. E. W. Johnson, L. Bulgarelli, L. Shen, A. Gayles, A. Shammout, S. Horng, T. J. Pollard, S. Hao, B. Moody, B. Gow, L.-w. H. Lehman, L. A. Celi, and R. G. Mark, "MIMIC-IV, a freely accessible electronic health record dataset," *Scientific Data*, vol. 10, no. 1, p. 1, 2023.

[3] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.

[4] A. Khwaja, "KDIGO clinical practice guidelines for acute kidney injury," *Nephron Clinical Practice*, vol. 120, no. 4, pp. c179–c184, 2012.

[5] E. W. Ely, S. K. Inouye, G. R. Bernard, S. Gordon, J. Francis, L. May, B. Truman, T. Speroff, S. Gautam, R. Margolin *et al.*, "Delirium in mechanically ventilated patients: validity and reliability of the confusion assessment method for the intensive care unit (CAM-ICU)," *JAMA*, vol. 286, no. 21, pp. 2703–2710, 2001.

[6] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.